

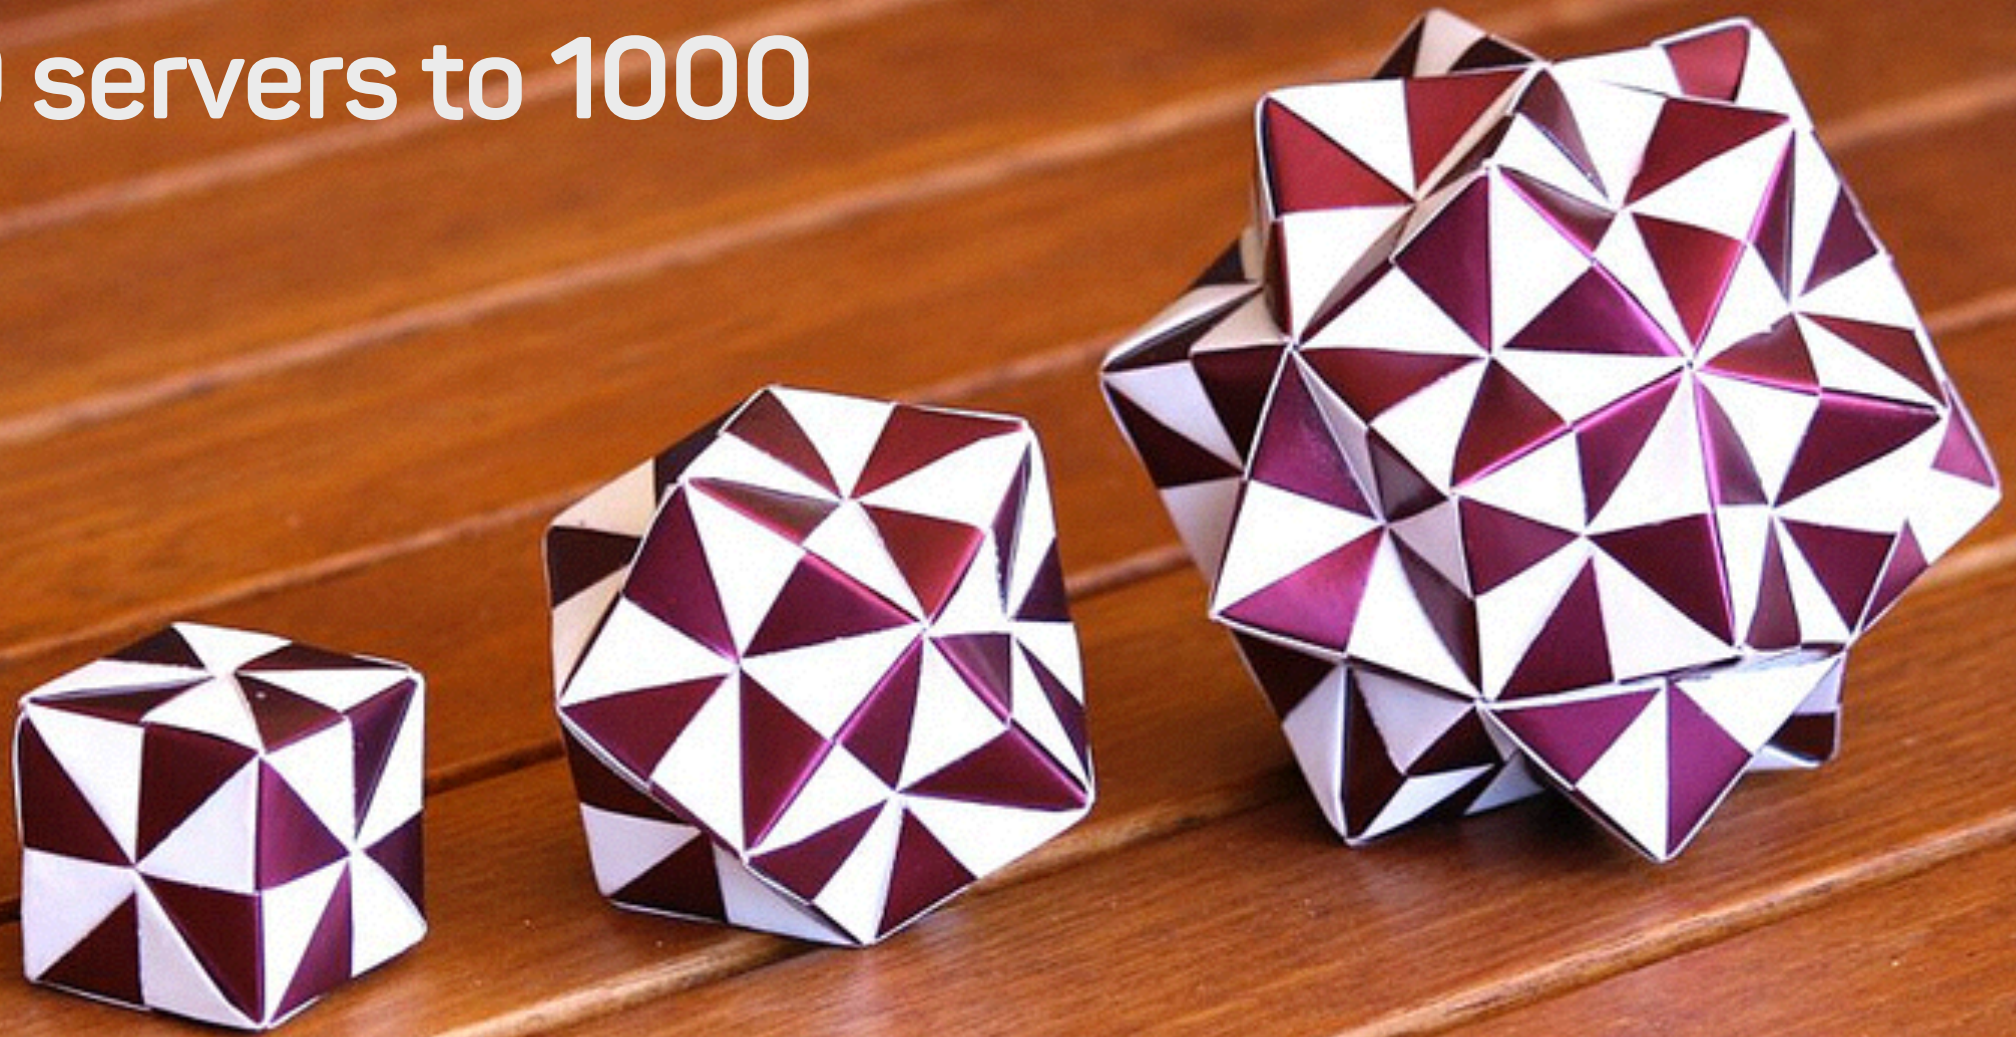
# Multi-Stage Ansible

---

Victor Volle  
Ansible Meetup München 2016



# From 10 servers to 1000



**DevOps Borat**  
@DEVOPS\_BORAT

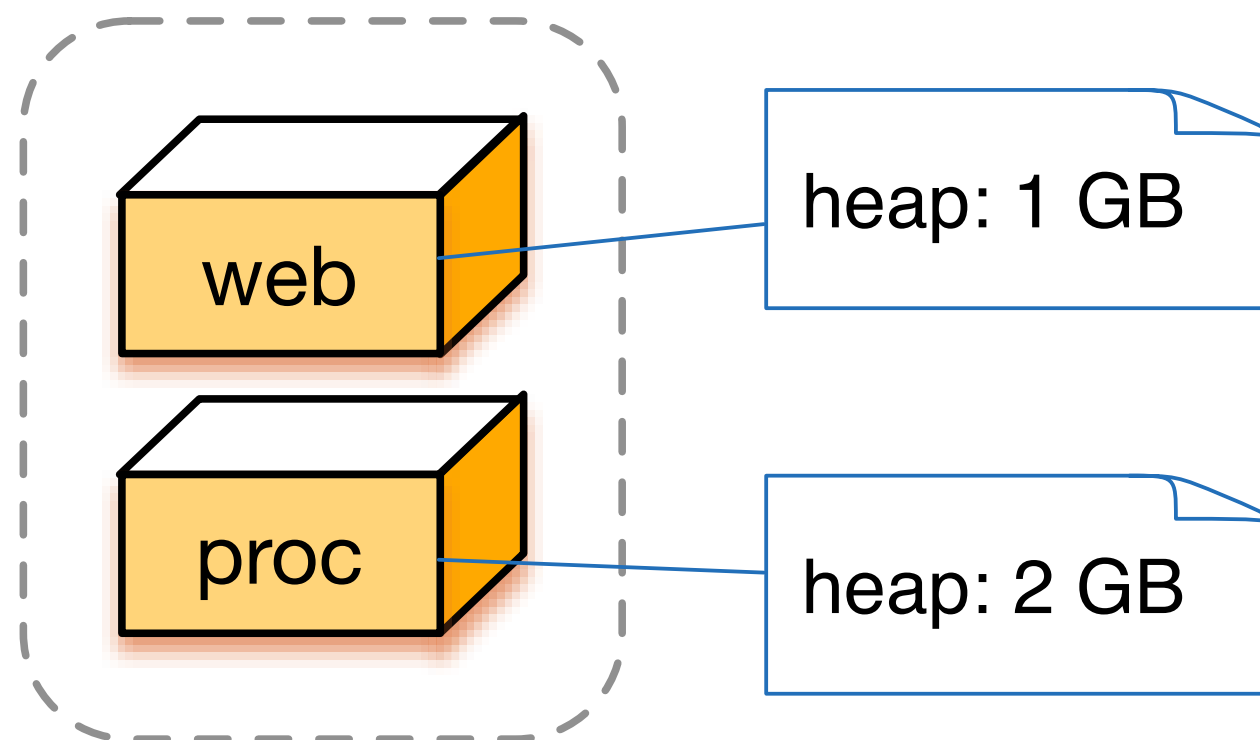
 Follow

To make error is human. To propagate error to all server in automatic way is [#devops](#).

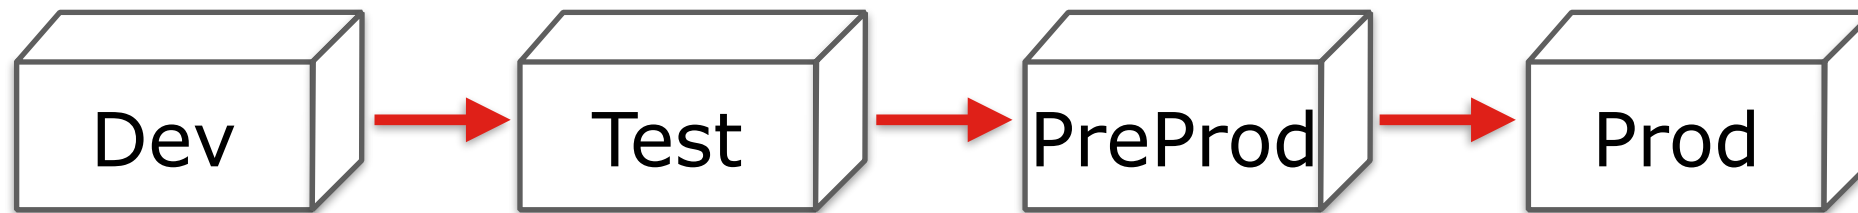
 Reply  Retweet  Favorite  More



## "configuration"



## "configuration"



web

**mango01**

heap=2GB

**mango02**

heap=2GB

**mango03-06**

heap=4GB

**mango07-10**

heap=4GB

proc

**dattel01**

heap=2GB

**dattel02**

heap=4GB

**dattel02,  
dattel03**

heap=8GB

**dattel07,  
dattel08**

heap=8GB

## multi-stage: possible solutions

1. group variables
2. multiple inventories
3. exploiting variable precedence
4. using "children"
5. creating your own vars plugin
6. talk to Brian Coca

## multi-stage: possible solutions

1. group variables
2. multiple inventories
3. exploiting variable precedence
4. using "children"
5. creating your own vars plugin
6. talk to Brian Coca

# group variables (inventory)

## # Inventory

```
[web]
mango[01-10]

[proc]
datte1[01-08]
```

```
[dev]
mango01
datte101
```

```
...
```

```
[prod]
mango[07-10]
datte1[07-08]
```

```
...
```

```
[web:vars]
heap=2G

[proc:vars]
heap=2G
```

```
[dev:vars]
heap=2G
```

group  
variables

it is not possible to set  
different values for  
"heap" for **web** and **proc**  
we need a "namespace"

# group variables (inventory)

## # Inventory

```
[web]
mango[01-10]

[proc]
datte1[01-08]
```

```
[dev]
mango01
datte101
```

...

```
[prod]
mango[07-10]
datte1[07-08]
```

...

```
[web:vars]
webheap=2G

[proc:vars]
procheap=2G
```

```
[dev:vars]
webheap=2G
procheap=4G
```

...

```
[prod:vars]
webheap=4G
procheap=8G
```

procheap

dev

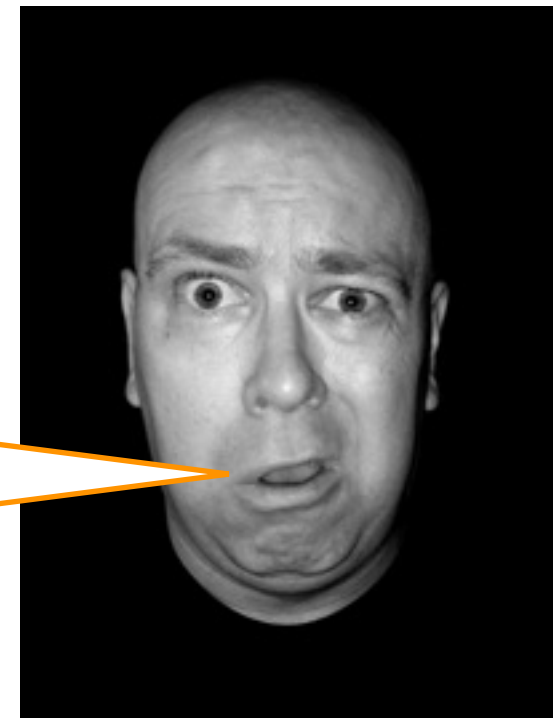
4G

prod

2G

groups are  
alphabetically  
sorted

"proc"  
< "prod"





# Default Workspace structure

group variables

```
workspace/  
├─ files/  
├─ group_vars/  
├─ host_vars/  
├─ roles/  
├─ templates/  
├─ ansible.cfg  
├─ ansible_inventory  
├─ java.yml  
...
```

groups are still  
sorted alphabetically

# Variables

» Avoid defining the variable "x" in 47 places and then ask the question "which x gets used". Why? Because that's not Ansible's Zen philosophy of doing things.

There is only one Empire State Building. One Mona Lisa, etc. Figure out where to define a variable, and don't make it complicated.

## multi-stage: possible solutions

1. group variables
- 2. multiple inventories**
3. exploiting variable precedence
4. using "children"
5. creating your own vars plugin
6. talk to Brian Coca



# inventories per environment

```
# dev_inventory
```

```
[web]  
mango01
```

```
[proc]  
dattel01
```

```
[dev:children]  
web  
proc
```

```
[dev:vars]  
webheap=2G  
procheap=4G
```

...

```
# prod_inventory
```

```
[web]  
mango02
```

```
[proc]  
dattel02
```

```
[prod:children]  
web  
proc
```

```
[prod:vars]  
webheap=4G  
procheap=8G
```

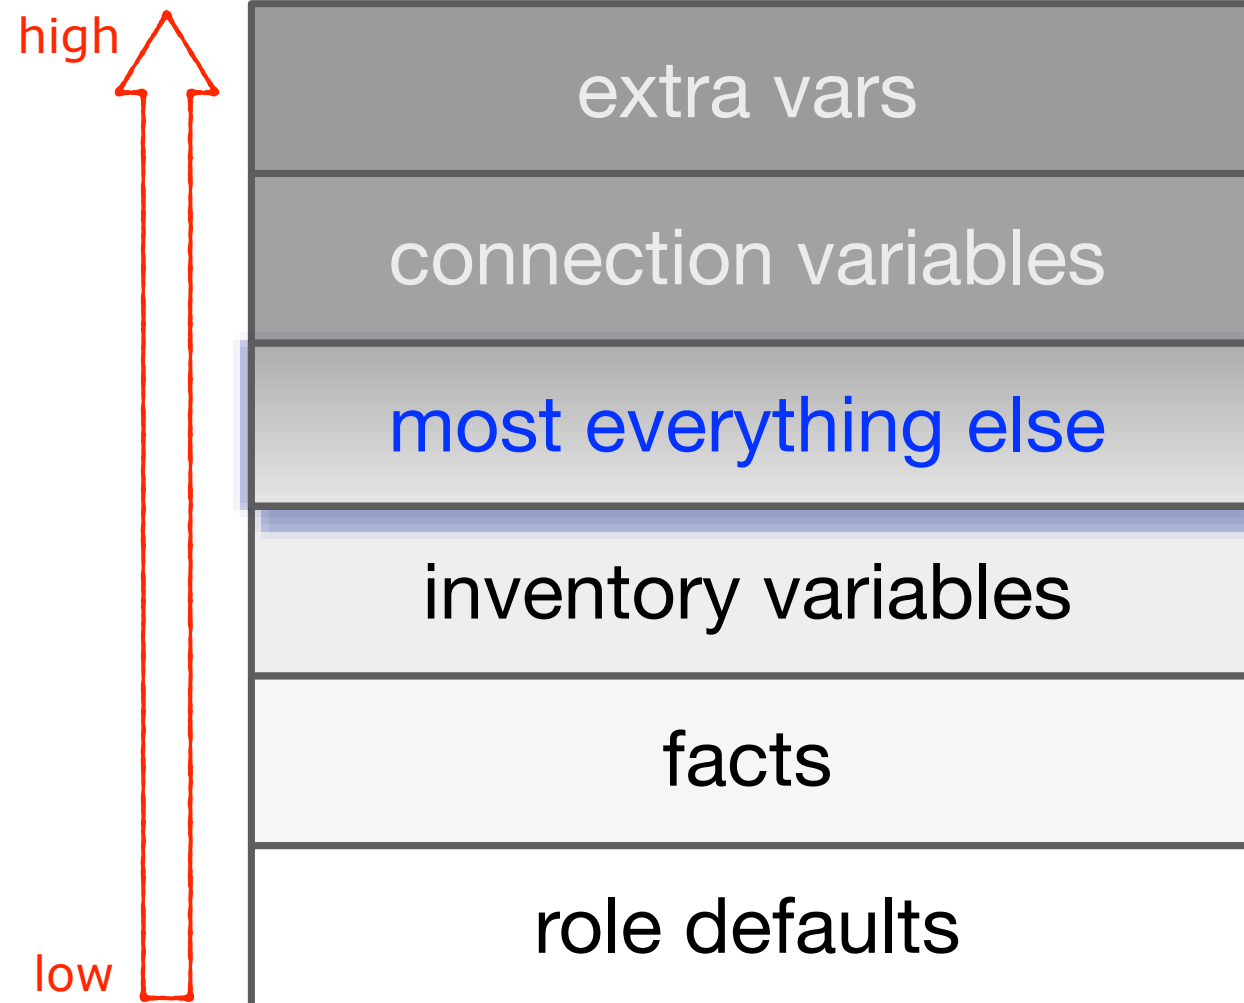
- **“good enough” in most cases! Use it as long as possible**
- *(you should place the variables in the group\_vars folder though)*
- you cannot invoke a playbook on *all* hosts — which you probably do not want anyway
- there is some duplication in the inventories, but not too much
- **get’s very unwieldy when you add another dimension like ‘location’**

see e.g. : <http://rosstuck.com/multistage-environments-with-ansible/>

## multi-stage: possible solutions

1. group variables
2. multiple inventories
- 3. exploiting variable precedence**
4. using "children"
5. creating your own vars plugin
6. talk to Brian Coca

# Variable Precedence



```
> ansible-playbook -e "heap=16G" ...
```

```
# Inventory
```

```
mango01 ansible_ssh_host=10.0.1.16
```

```
# Inventory
```

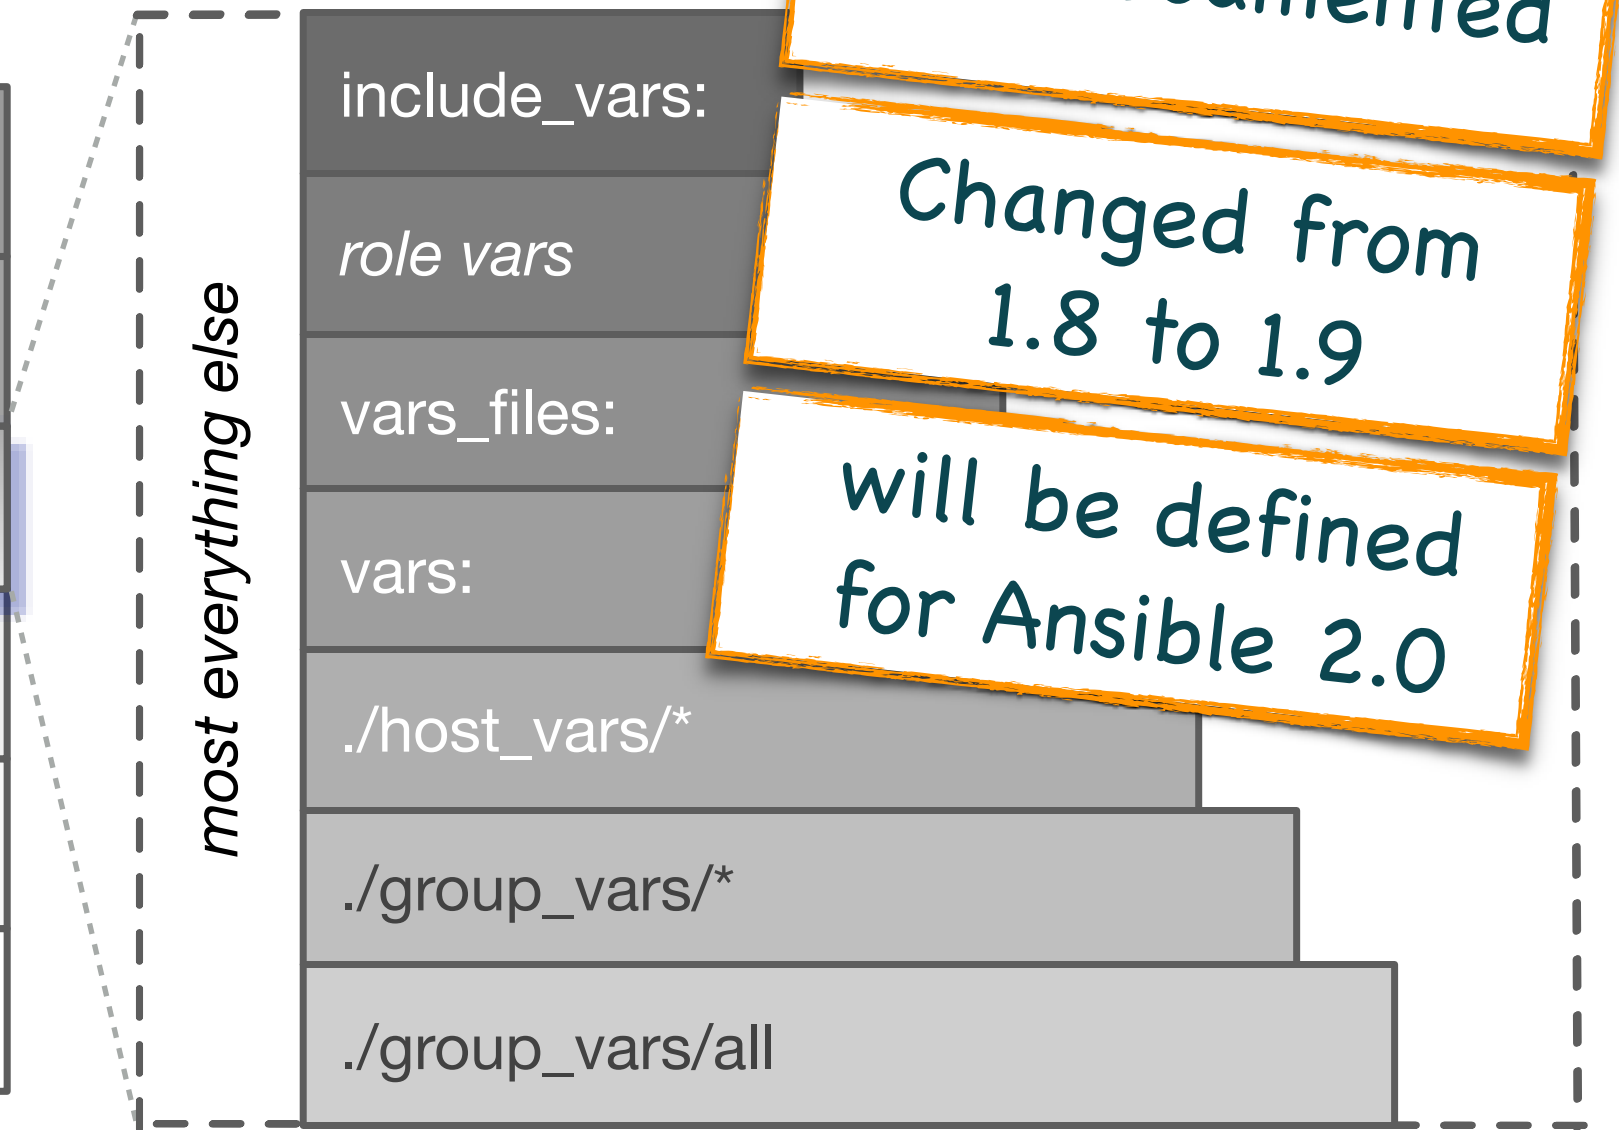
```
...
```

```
mango10 heap=8G
```

```
[web:vars]  
heap=4G
```

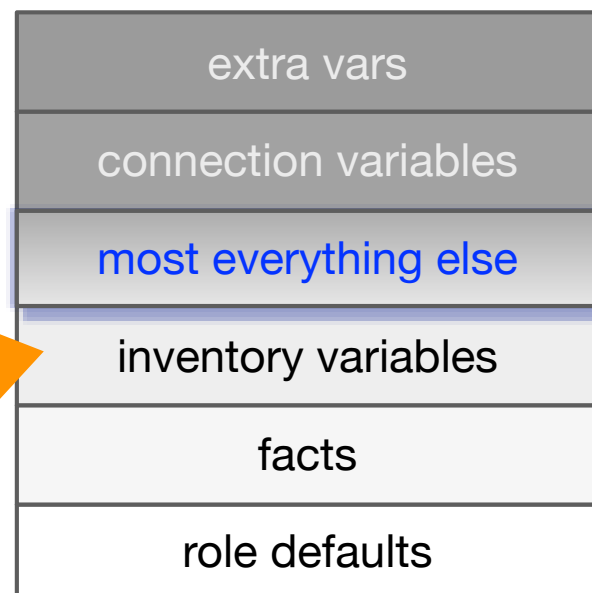


# Variable Precedence

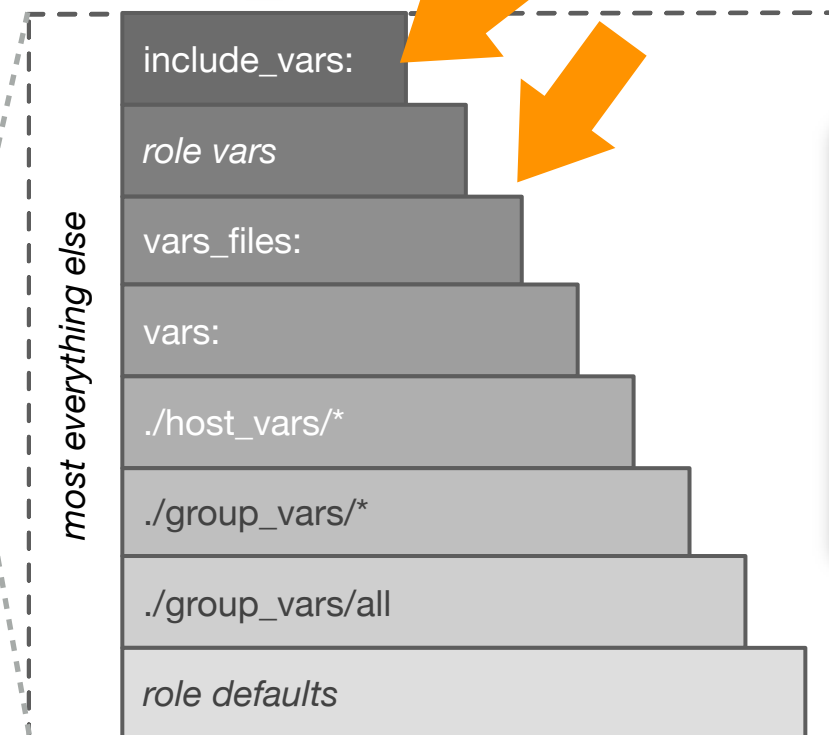


## variable precedence

Inventory per Environment  
Problem: inventory variables have a low priority



Use *include\_vars* task or *vars\_files*  
(That's what we are currently doing)  
Problem: you have to add that to every Playbook



Write your own vars plugin  
(That's what I am currently planning)

# Variables per stage/envrionment

```
# varprecedence.yml
---
- name: Check Var precedence
  hosts: "{{ lookup('env','STAGE') }}:&web"

  vars_files:
    - "{{ lookup('env','STAGE') }}_vars.yml"

  tasks:
    - include_vars: include_vars{{ lookup('env','STAGE') }}.yml
    ...
```

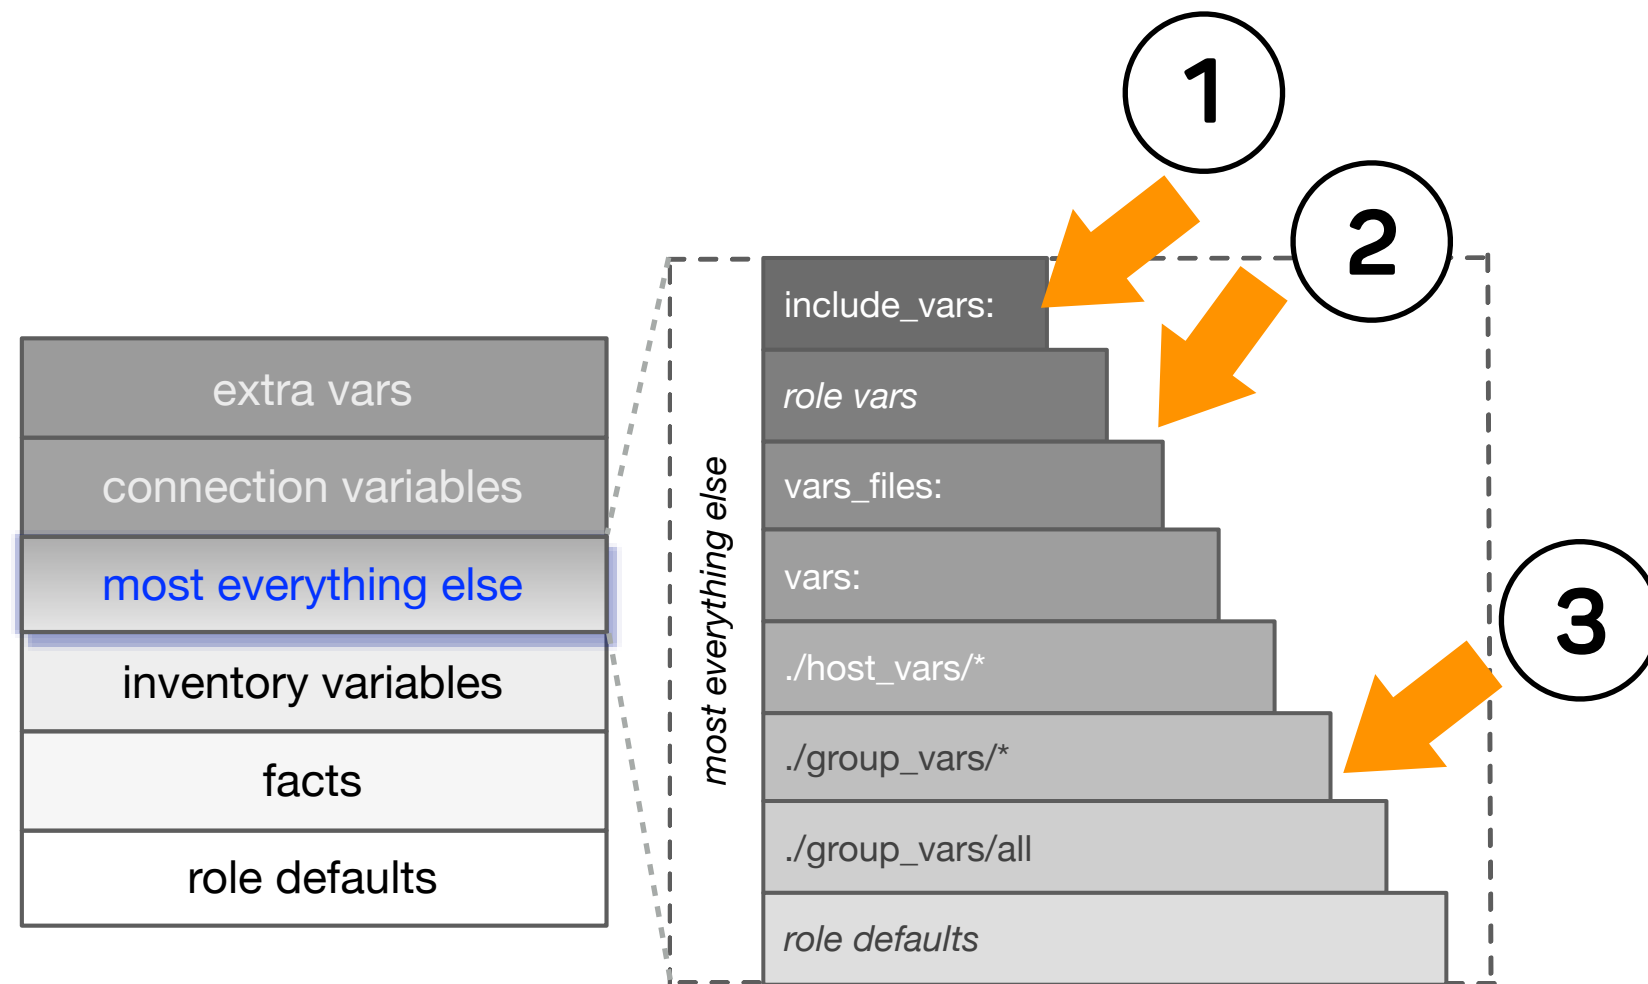
servers, which are in  
group "web" and in  
group "{{STAGE}}"

(1) load variables from  
environment/stage  
specific file

(2) load variables from  
environment/stage  
specific file



## “exploiting variable precedence”



## multi-stage: possible solutions

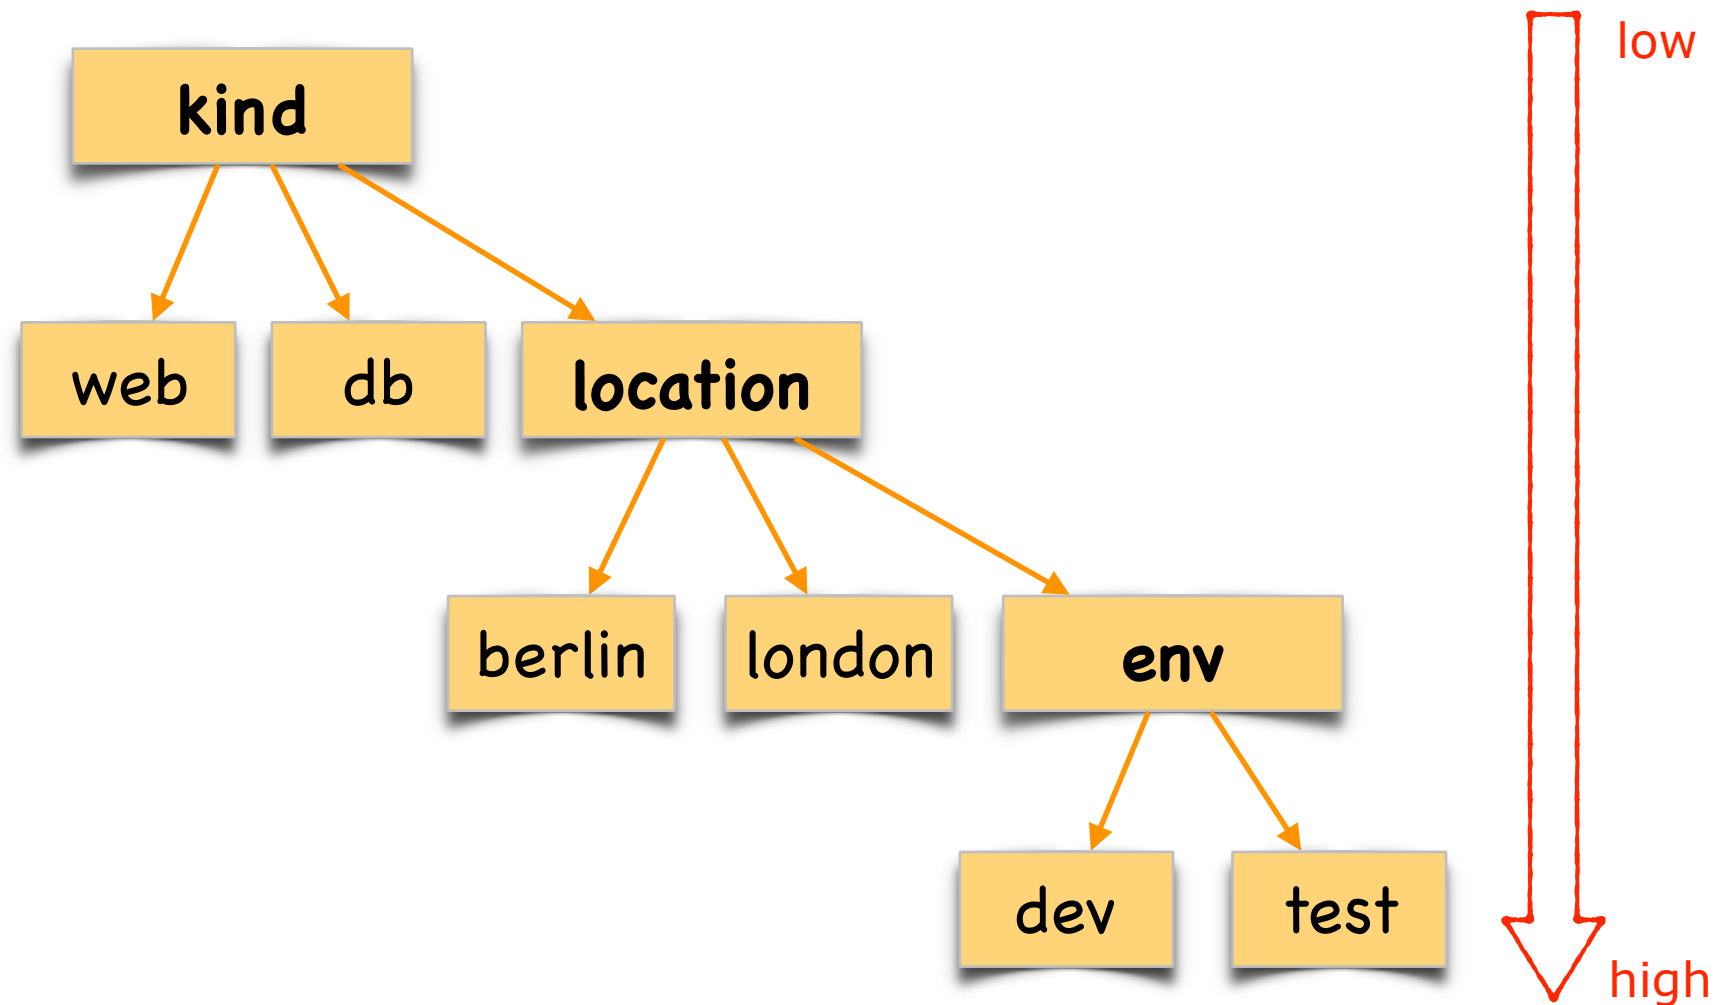
1. group variables
2. multiple inventories
3. exploiting variable precedence
- 4. using "children"**
5. creating your own vars plugin
6. talk to Brian Coca

# children have higher priority than their parent

```
[env:children]  
dev  
test
```

```
[location:children]  
berlin  
env  
london
```

```
[kind:children]  
db  
location  
web
```



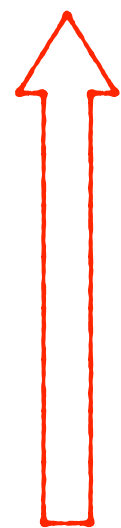


## multi-stage: possible solutions

1. group variables
2. multiple inventories
3. exploiting variable precedence
4. using "children"
- 5. creating your own vars plugin**
6. talk to Brian Coca

## Puppet Hiera: a simple Hierarchical Database

`:hierarchy:`

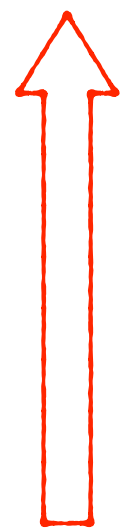


- `- "locations/%{::location}"`
- `- "envs/%{::stage}"`
- `- "roles/%{::role}"`
- `- common`

a puppet "role" is  
roughly equivalent to  
an Ansible "group"

## Puppet Hiera: a simple Hierarchical Database

`:hierarchy:`



- `- "locations/%{::location}"`
- `- "envs/%{::stage}"`
- `- "roles/%{::role}"`
- `- common`

```
├─ locations/  
|   ├─ ams.yaml  
|   ├─ asia.yaml  
|   ├─ berlin.yaml  
|   ├─ praha.yaml  
|   └─ timb.yaml  
├─ envs/  
|   ├─ dev.yaml  
|   ├─ preprod.yaml  
|   ├─ prod.yaml  
|   └─ test.yaml  
└─ roles/  
    ├─ proc.yaml  
    └─ web.yaml
```

<https://github.com/kontrafiktio/ansible-plugins>

```
class Hierme(object):

    def __init__(self, hieme_file_name):
        defaults = {"hash_behavior": default_hash_behaviour}
        self.data = {}
        self.hieme_file_name = hieme_file_name
        with open(hieme_file_name, 'r') as hierme_file:
            docs = yaml.load_all(hierme_file)
            doc = deep_merge(defaults, docs.next())
            self.hash_behavior = doc["hash_behavior"]
            self.hierarchy = doc["hierarchy"]
            if type(self.hierarchy).__name__ == "str":
                self.hierarchy = [self.hierarchy]
            self.backend = YamlBackend(doc["datadir"])

    def run(self, host, vault_password=None):
        for layer in self.hierarchy:
            new_data = self.backend.read(layer, wrap=True)
            if self.hash_behavior == HASH_BEHAVIOUR_MERGE:
                self.data = deep_merge(new_data, self.data)
            else:
                for key in new_data:
                    if not key in self.data:
                        self.data[key] = new_data[key]
```

nothing to see there,  
yet



## multi-stage: possible solutions

1. group variables
2. multiple inventories
3. exploiting variable precedence
4. using "children"
5. creating your own vars plugin
- 6. talk to Brian Coca**

# Februar 2015: Mailing list

Ansible Project ›

multi stage and variables

10 posts by 6 authors 



me (Victor Volle [change](#))

2/1/15



Hi!

I tried to read up on everything I could find w.r.t. multi stage environments, but I haven't found a solution that fits my needs yet.

Let's say, I have 4 "environments": dev, test, stage, production  
And 2 "kinds" of servers: web, db

To simplify things, I only have one variable, the size of the "ram".  
In "dev", "test" and "stage" the ram for "web" should be 1GB, and 4GB for "prod" (artificial example to show the issue!)  
And in "dev", "test" and "stage" the ram for "db" should be 2GB, and 8GB for "prod":

	dev	test	stage	prod
web	+-----+   (1GB)   +-----+	+-----+   (1GB)   +-----+	+-----+   (1GB)   +-----+	+-----+   (4GB)   +-----+
db	+-----+   (2GB)   +-----+	+-----+   (2GB)   +-----+	+-----+   (2GB)   +-----+	+-----+   (8GB)   +-----+

# August 2015: Issue 12156

## Feature Idea: Extendable Variable Manager #12156

 **Closed**

**kontrafiktion** opened this issue on Aug 29, 2015 · 2 comments



**kontrafiktion** commented on Aug 29, 2015

I would like to be able to extend/replace the Ansible (2.0) VariableManager, so that I would be able to define a custom hierarchy of groups:

group category "stage" contains groups: "dev", "test", "preprod" and "prod"  
other groups: "web", "db"

whatever is defined as group vars for "web" or "db" should be overridden by anything in the group category "stage".

# February 2016: #ansiblefest

group precedence for variables #14556

 **Open** kontrafiktion opened this issue 2 days ago · 3 comments



kontrafiktion commented 2 days ago

**Issue Type:**

- Feature Idea

**Ansible Version:**

Ansible 2.0

**Ansible Configuration:**

n/a

**Environment:**

n/a

**Summary:**

Today at #ansiblefest I sat down with @bcoca about variable/group precedence.

That's already  
on my todo list

## Who me?

**Dr. Victor Volle**

*IT-Something*



**"Senior IT Consultant"**

**2015- codecentric**

**"Architect"**

**2008- Senacor Technologies AG**

**"Chief Architect JEE"**

**2004- ING DiBa**

**Developer, Architect, Head of ...**

**1996- develop group, Erlangen**



